



Chapter 5: Concurrency Control Techniques

Adama Science and Technology University
School of Electrical Engineering and Computing
Department of CSE
CSEg 2208: Database Systems
(2022)

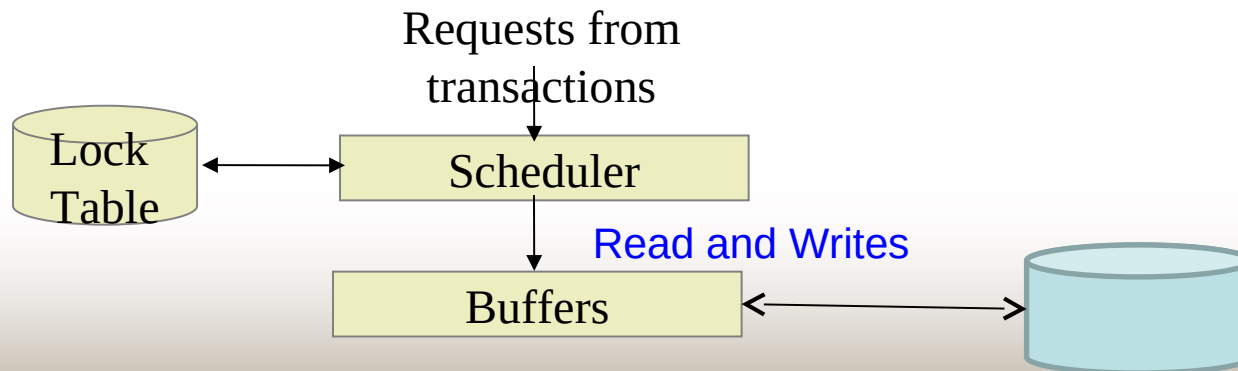
Outline

- ❖ Databases Concurrency Control
- ❖ Purpose of Concurrency Control
- ❖ Concurrency Control Techniques:
 - ❖ Locking
 - ❖ Timestamp
 - ❖ Optimistic
 - ❖ Multiversion
 - ❖ Lock Granularity

Database Concurrency Control

Transaction Processor is divided into:


- A **concurrency-control manager/scheduler**, responsible for assuring *isolation of transactions*.
- A **logging and recovery manager**, responsible for the *durability of transactions*.
- The scheduler (**concurrency-control manager**) must assure that the individual actions of *multiple transactions* are executed in such an order that the net effect is the same as if the transactions had in fact *executed one-at-a-time*.



Purpose of Concurrency Control

Example:

Bank database: 3 Accounts



Property: $A + B + C = 1500$

Money does not *leave the system*

Purpose of Concurrency Control

- To *enforce Isolation* (through mutual exclusion) among conflicting transactions.
- To *preserve database consistency* through *consistency preserving* execution of transactions.
- To resolve *read-write* and *write-write* conflicts.
- A *typical scheduler* does its (concurrency control) work by maintaining *locks on certain pieces* of the database.
- These locks *prevent two transactions* from accessing the same piece of data at the same time. **Example:**
 - In concurrent execution environment if T_1 conflicts with T_2 over a data item A , then the *existing concurrency control* decides if T_1 or T_2 should get the A and if the other transaction is *rolled-back or waits*.

Purpose of Concurrency Control

Example

[Transaction T1: Transfer 100 from A to B]

Read (A, t)

$t = t - 100$

Write (A, t)

Read (B, t)

$t = t + 100$

Write (B, t)

[Transaction T2: Transfer 100 from A to C]

Read (A, s)

$s = s - 100$

Write (A, s)

Read (C, s)

$s = s + 100$

Write (C, s)

Purpose of Concurrency Control

Transaction T_1	Transaction T_2	A	B	C
Read (A, t)		500	500	500
$t = t - 100$				
	Read (A, s)			
	$s = s - 100$			
	Write (A, s)	400	500	500
Write (A, t)		400	500	500
Read (B, t)				
$t = t + 100$				
Write (B, t)		400	600	500
	Read (C, s)			
	$s = s + 100$			
	Write (C, s)	400	600	600

$$400 + 600 + 600 = 1600$$

Schedule

Purpose of Concurrency Control

Transaction T1	Transaction T2	A	B	C
Read (A, t)		500	500	500
$t = t - 100$				
Write (A, t)		400	500	500
	Read (A, s)			
	$s = s - 100$			
	Write (A, s)	300	500	500
Read (B, t)				
$t = t + 100$				
Write (B, t)		300	600	500
	Read (C, s)			
	$s = s + 100$			
	Write (C, s)	300	600	600

$$300 + 600 + 600 = 1500$$

So What ?

Alternative Schedule

Concurrency Control Techniques

- **Basic concurrency control techniques:**
 - **Locking,**
 - **Timestamping**
 - **Optimistic methods**
- The First two are conservative approaches: delay transactions in case they *conflict with other transactions*.
- Optimistic methods assume conflict is *rare and only check for conflicts at commit*.
- **Locking:**
 - Lock is a *variable* associated with a *data item* that *describes the status of the data item* with respect to the possible operations that can be applied to it.

Concurrency Control Techniques

- Generally, a transaction must claim a *shared (read)* or *exclusive (write)* lock on a data item before *read or write*.
- *Lock prevents* another transaction from *modifying item or even reading it*, in the case of a write lock.
- **Locking is an operation which secures :**
 - (a) permission to *Read*,
 - (b) permission to *Write a data item* for a transaction.
 - **Example:**
 - Lock (X). Data item X is locked in behalf of the *requesting transaction*.
- **Unlocking** is an operation *which removes* these permissions from the data item.
 - **Example:**
 - Unlock (X): Data item X is made available to *all other transactions*.
- Lock and Unlock are **Atomic operations**.

Concurrency Control Techniques

■ Two locks modes:

- (a) shared (read) (b) exclusive (write).
- *Shared mode: shared lock (X)*
 - *More than one transaction* can apply share lock on X for reading its value but *no write lock* can be applied on X by any other transaction.
- *Exclusive mode: Write lock (X)*
 - *Only one write lock* on X can exist at any time and *no shared lock* can be applied by any other transaction on X.

■ Conflict matrix:

	Read	Write
Read	Y	N
Write	N	N

Concurrency Control Techniques

- Lock Manager:

- Managing *locks* on data items.

- Lock table:

- Lock manager uses it to store the *identify of transaction locking a data item*, the *data item*, *lock mode* and *pointer* to the next data item locked.
- One simple way to implement a lock table is through *linked list*.

Transaction ID	Data item id	lock mode	Ptr to next data item
T1	X1	Read	Next

- Database requires that all transactions should be *well-formed*.

- A transaction is *well-formed* if:

- It must *lock the data item* before it *reads* or *writes* to it.
- It must *not lock* an already locked data items and it must not try to unlock a free data item.

Concurrency Control Techniques

Locking - Basic Rules:

- It has two operations : **Lock_item(X)** and **unLock_item(X)**.
- A transaction request access to an item X by first issuing a **lock_Item(x)** operation .
- If $\text{lock}(x)=1$, the transaction is **forced to wait**.
- If $\text{lock}(X)=0$; it **is set to 1** and the transaction is allowed to access x.
- When a transaction finished operation on X it issues an **Unlock_item operation** which **set lock(x)** to 0 so that X may be accessed by another transaction
- If transaction has **shared lock on item**, can read but not update item.
- If transaction has **exclusive lock on item**, can both read and update item.
- Reads cannot conflict, so **more than one transaction** can **hold shared locks** simultaneously on same item.
- Exclusive lock gives transaction **exclusive access** to that item.

Concurrency Control Techniques

- The following code performs the read operation: *read_lock(X)*:

```
B: if LOCK (X) = "unlocked" then
    begin
        LOCK (X) ← "read-locked";
        no_of_reads (X) ← 1;
    end
else if LOCK (X) ← "read-locked" then
    no_of_reads (X) ← no_of_reads (X) + 1
else begin
    wait (until LOCK (X) = "unlocked" and
        the lock manager wakes up the transaction);
    go to B
end;
```

Concurrency Control Techniques

- The following code performs the write lock operation: *write_lock(X)*:

B: if LOCK (X) = “unlocked” then

 LOCK (X) \leftarrow “write-locked”;

else

 wait (until LOCK(X) = “unlocked”

 and the lock manager wakes up the transaction);

goto **B**

end;

Concurrency Control Techniques

- **Lock conversion**
 - *Lock upgrade*: existing read lock to *write lock*:
if T_i has a read-lock (X) and T_j has no read-lock (X) ($i \neq j$) then
convert read-lock (X) to write-lock (X)
else
force T_i to wait until T_j unlocks X
 - *Lock downgrade*: existing write lock to *read lock*:
 T_i has a write-lock (X) (*no transaction can have any lock on X*) convert
write-lock (X) to read-lock (X)
- Using such locks in the transaction do not guarantee serializability of schedule on its own: Example

Concurrency Control Techniques

Schedule:

T1

T2

Write Lock(X)
Read (X)
X=X+100
Write(X)
Unlock(X)

Write Lock(X)

Read (X)
X= X*1.1
Write(X)
Unlock(X)

Write Lock(Y)

Read (Y)
Y= Y*1.1
Write(Y)
Unlock(Y)
Commit

write_lock(Y)
read(Y)
Y= Y-100
write(Y)
unlock(Y)
commit

**Example : Incorrect
Locking Schedule**

Concurrency Control Techniques

- If at start, $X = 100$, $Y = 400$, result should be:
 - $X = 220$, $Y = 330$, if T_1 executes before T_2 , or
 - $X = 210$, $Y = 340$, if T_2 executes before T_1
- However, result gives $X = 220$ and $Y = 340$.
- S is not a serializable schedule. **Why?**
- Problem is that transactions release locks too early (soon), resulting in loss of *total isolation* and *atomicity*.
- To guarantee serializability, we need an *additional protocol* concerning the *positioning of lock* and *unlock operations* in every transaction.

Concurrency Control Techniques


2.2 Two-Phase Locking Techniques: The algorithm

- Transaction follows 2PL protocol if all *locking* operations *precede first unlock operation* in the transaction.
- Every transaction can be divided into *Two Phases: Locking (Growing) & Unlocking (Shrinking)*
 - Locking (Growing) Phase:*
 - A transaction applies locks (read or write) on *desired data items* one at a time.
 - Acquires *all locks but cannot release any locks*.
 - Unlocking (Shrinking) Phase:*
 - A transaction unlocks its *locked data items one at a time*.
 - Releases locks but *cannot acquire any new locks*.

Concurrency Control Techniques

Requirement:

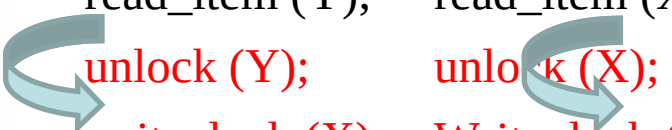
- For a transaction *these two phases* must be *mutually exclusively*, that is, *during locking phase unlocking phase* must not start and during *unlocking* phase locking phase must not begin.



Concurrency Control Techniques

Example

T1 T2



read_lock (Y);	read_lock (X);
read_item (Y);	read_item (X);
unlock (Y);	unlock (X);
write_lock (X);	Write_lock (Y);
read_item (X);	read_item (Y);
X:=X+Y;	Y:=X+Y;
write_item (X);	write_item (Y);
unlock (X);	unlock (Y);

Result

Initial values: X=20; Y=30

Result of serial execution

T1 followed by T2

X=50, Y=80.

Result of serial execution

T2 followed by T1

X=70, Y=50

Concurrency Control Techniques

T'1 T'2

read_lock (Y);	read_lock (X);	
read_item (Y);	read_item (X);	
write_lock (X);	Write_lock (Y);	
unlock (Y);	unlock (X);	.
read_item (X);	read_item (Y);	
X:=X+Y;	Y:=X+Y;	
write_item (X);	write_item (Y);	
unlock (X);	unlock (Y);	

- T'1 and T'2 follow two-phase policy but they are subject to *deadlock*, which must be dealt with

- If every transaction in the schedule follows the 2PL protocol , the schedule is guaranteed to be *serializable*.
- **Limitation** -It may limit the number of concurrency that can occur in the schedule .
How?
- **Remark:** The use of this locking can cause two additional problems
 - **Deadlock and starvation**

Concurrency Control Techniques

Dealing with Deadlock and Starvation

- **Deadlock**

- It is a state that may result when *two or more transaction* are each waiting for locks held by the other to be released.
- **Example :**

T1

read_lock (Y);
read_item (Y);

write_lock (X);

T2

read_lock (X);
read_item (X);

write_lock (Y);

Concurrency Control Techniques

- T1 is in the *waiting queue* for X which is locked by T2.
- T2 is on the *waiting queue* for Y which is locked by T1.
- No transaction can *continue until the other transaction completes*.
- T1 and T2 did follow two-phase policy but they are deadlock.
- So the DBMS must either *prevent* or *detect* and *resolve* such deadlock situations

There are possible solutions : Deadlock prevention, deadlock detection and avoidance ,and lock timeouts

i. Deadlock prevention protocol: two possibilities

■ The conservative two-phase locking

- A transaction locks *all data items it refers* to before it begins execution.
- This way of locking *prevents deadlock* since a transaction never waits for a data item.

07/14/22 Limitation : *It restricts concurrency*

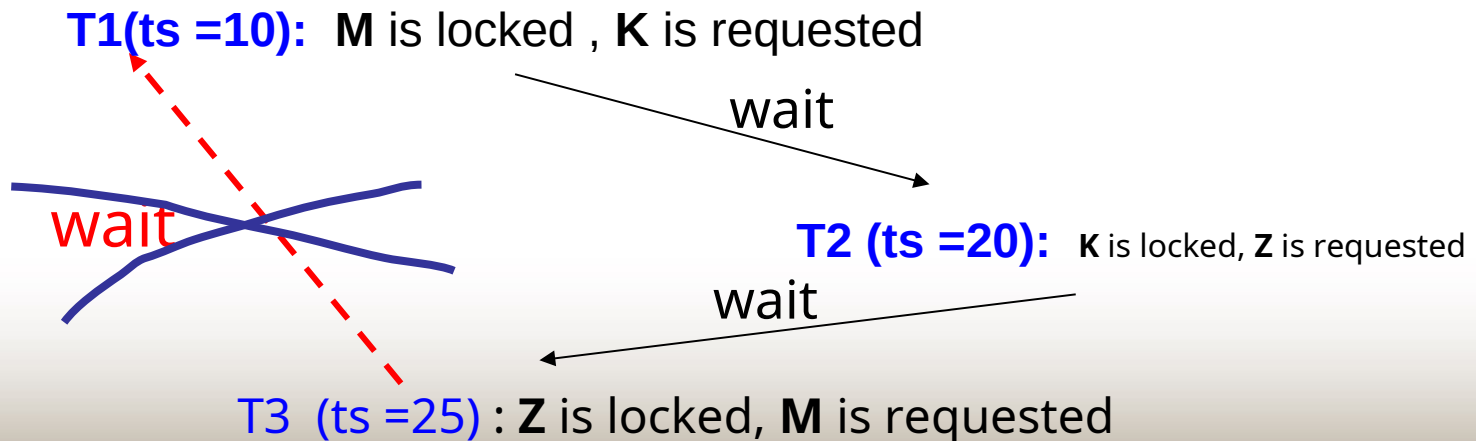
Concurrency Control Techniques

- Transaction Timestamp($TS(T)$)
 - We can prevent deadlocks by *giving each transaction a priority* and ensuring that *lower priority transactions* are not allowed to wait for higher priority transactions (or vice versa).
 - One way to *assign priorities* is to give each transaction a *timestamp* when it starts up.
 - it is a *unique identifier given to each transaction* based on time in which it is started. i.e if T_1 starts before T_2 , $TS(T_1) < TS(T_2)$
 - The *lower the timestamp*, the *higher the transaction's priority*, that is, the oldest transaction has the highest priority.
 - **If a transaction T_i requests a lock and transaction T_j holds a conflicting lock, the lock manager can use one of the following two policies: Wait-die & Would-wait**

Concurrency Control Techniques

- **Wait-die**


- If T_i has *higher priority*, it is allowed to wait; otherwise it is aborted.
- An older transaction is allowed to wait on a younger transaction.
- A *younger transaction requesting* an item held by an older transaction is *aborted*.
- If $TS(T_i) < TS(T_j)$, then (T_i older than T_j) T_i is allowed to *wait*.
- Otherwise (T_i younger than T_j) *Abort* T_i (T_i dies) and restart it later with the same timestamp.



Concurrency Control Techniques

- **Wound-wait**

- The opposite of wait-die.
- If T_i has *higher priority*, abort T_j ; otherwise T_i waits.
- A younger transaction is allowed to wait on an older one
- An older transaction requesting an item held by a younger transaction preempts the younger transaction by aborting it.
- If $TS(T_i) < TS(T_j)$, then (T_i older than T_j), **Abort** T_j (T_i wounds T_j) and restart T_j later with the same timestamp.
- Otherwise (T_i younger than T_j) T_i is allowed to *wait*.





Concurrency Control Techniques

■ Remark:

- Both methods ended by aborting *the younger of the two transaction* that may be involved in the deadlock.
- Limitation:
 - Both techniques may cause *some transaction* to be *aborted* and *restarted* needlessly.

ii. Deadlock Detection and resolution

- In this approach, deadlocks are *allowed to happen*.
- The scheduler maintains a *wait-for-graph for detecting cycle*.
- When a chain like: T_i waits for T_j waits for T_k waits for T_i or T_j occurs, then this creates a cycle.
- When the system is in the *state of deadlock*, some of the transaction should be aborted by selection(victim) and rolled-back.
- This can be done by aborting those transaction: *that have made the least work, the one with the lowest locks, and that have the least # of abortion and so on*
- Example:



Concurrency Control Techniques

iii. Timeouts

- It uses the *period of time* that several transaction have been waiting to lock items.
- It has *lower overhead cost* and it is *simple*.
- If the *transaction wait for a longer time* than the predefined time out period, the system assume that may be *deadlocked* and *aborted* it.

■ Starvation

- Starvation occurs when a particular transaction *consistently waits* or *restarted* and never gets a chance to proceed further while other transaction continue normally.
- This may occur , if the *waiting method* for item locking:
 - Gave priority for some transaction over others.
 - Problem in Victim selection algorithm- it is possible that the same transaction may consistently be selected as victim and rolled-back .example In **Wound-Wait**
- **Solution**
 - FIFO,
 - Allow for transaction that *wait for a longer time*,
 - Give *higher priority* for transaction that have been *aborted for many times*.

Concurrency Control Techniques

Timestamp based concurrency control algorithm

- **Timestamp**

- In lock based concurrency control , conflicting actions of different transactions are ordered by the order in which locks are obtained.
- But here, Timestamp values are assigned based on time in which the transaction are submitted to the system using the current date & time of the system
- A monotonically increasing variable (integer) indicating the age of an operation or a transaction.
- A larger timestamp value indicates a more recent event or operation.
- Timestamp based algorithm uses timestamp to serialize the execution of concurrent transactions.
- It **doesn't use lock**, thus **deadlock cannot** be occurred
- In the timestamp ordering, conflicting operation in the schedule shouldn't violate serializable ordering
- This can be achieved by associating timestamp value (TS) to each database item which is denoted as follow:

Concurrency Control Techniques

- a) **Read_Ts(x)**: the read timestamp of x – this is the largest time among all the time stamps of transactions that have successfully read item X.
- b) **Write_TS(X)**: the largest of all the timestamps of transactions that have successfully written item X.
- The concurrency control algorithm check whether conflict operation violate the timestamp ordering in the following manner: **three options**
 - i. **Basic Timestamp Ordering**
 - Transaction T issues a **write_item(X)** operation:
 - If $\text{read_TS}(X) > \text{TS}(T)$ or if $\text{write_TS}(X) > \text{TS}(T)$, then a younger transaction has already read/write the values of the data item x before T had a chance to write X . so abort and roll-back T and **restarted with a new, larger timestamp**. Why is with new timestamp?, is there a difference b/n this timestamp protocol and the 2PL for dead lock prevention?
 - If the condition above does not exist, then execute **write_item(X)** of T and set **write_TS(X)** to **TS(T)**.

Concurrency Control Techniques

❑ Transaction T issues a read_item(X) operation:

- If $\text{write_TS}(X) > \text{TS}(T)$, then a younger transaction has already written to the data item, so abort and roll-back T and reject the operation.
 - If $\text{write_TS}(X) \leq \text{TS}(T)$, then execute read_item(X) of T and set read_TS(X) to the larger of TS(T) and the current read_TS(X)
- **Limitation:** cyclic restart/starvation may occur when a transaction is continuously aborted and restarted

Concurrency Control Techniques

4. Multiversion Concurrency Control Techniques

- This approach maintains a *number of versions of a data item* and *allocates the right version* to a read operation of a transaction.
- Thus unlike other mechanisms a *read operation* in this mechanism is *never rejected*.
- This algorithm uses the concept of *view serializability* than *conflict serializability*.
- **Side effect:**
 - Significantly more storage (RAM and disk) is required to *maintain multiple versions*. To check unlimited growth of versions, a garbage collection is run when some criteria is satisfied.
- Two schemes : **based on time stamped ordering & 2PL**

Concurrency Control Techniques

i. Multiversion technique based on timestamp ordering

- Assume X_1, X_2, \dots, X_n are the version of a data item X created by a *write operation of transactions*.
- With each X_i a read_TS (read timestamp) and a write_TS (write timestamp) are associated.)
 - **read_TS(X_i)**: The read timestamp of X_i is the largest of all the timestamps of transactions that have successfully read version X_i .
 - **write_TS(X_i)**: The write timestamp of X_i that wrote the value of version X_i .
 - A new version of X_i is created only by a *write operation*.

Concurrency Control Techniques

- To ensure serializability, the following two rules are used:
 - i. If transaction **T** issues **write_item (X)** and version i of X has the highest $\text{write_TS}(X_i)$ of all versions of X that is also less than or equal to $\text{TS}(T)$, and $\text{read_TS}(X_i) > \text{TS}(T)$, then abort and roll-back T ; otherwise create a new version X_i and *set $\text{read_TS}(X) = \text{write_TS}(X_i) = \text{TS}(T)$.*
 - ii. If transaction **T** issues **read_item (X)**, find the version i of X that has the highest $\text{write_TS}(X_i)$ of all versions of X that is also less than or equal to $\text{TS}(T)$, then return the value of X_i to T , and *set the value of $\text{read_TS}(X_i)$ to the largest of $\text{TS}(T)$ and the current $\text{read_TS}(X_i)$.*
- Note that: *Rule two indicates that read request will never be rejected*

Concurrency Control Techniques

- **ii. Multiversion Two-Phase Locking Using Certify Lock**
 - Allow a transaction **T'** to read a data item **X** while it is write locked by a conflicting transaction **T**.
 - This is accomplished by maintaining two versions of each data item **X** where one version must always have been written by some committed transaction. This means a write operation always creates a new version of **X**.

Steps

1. **X** is the committed version of a data item.
2. **T** creates a second version **X'** after obtaining a write lock on **X**.
3. Other transactions continue to read **X**.
4. **T** is ready to commit so it obtains a certify lock on **X'**.
5. The committed version **X** becomes **X'**.
6. **T** releases its certify lock on **X'**, which is **X** now.

Concurrency Control Techniques

Compatibility tables for

read/write locking scheme

	Read	Write
Read	yes	no
Write	no	no

read/write/certify locking scheme

	Read	Write	Certify
Read	Yes	Yes	No
Write	Yes	No	No
Certify	No	No	No

Note:

- In multiversion 2PL read and write operations from conflicting transactions can be processed concurrently.
- This improves concurrency but it may delay transaction commit because of obtaining certify locks on all its writes.
- It avoids cascading abort but like strict two phase locking scheme conflicting transactions **may get deadlocked**.

Concurrency Control Techniques

Validation (Optimistic) Concurrency Control Schemes

- This technique allow transaction to *proceed asynchronously* and only at the time of commit, serializability is checked & transactions are aborted in case of non-serializable schedules.
- Good if there is *little interference* among transaction.
- It has three phases: ***Read, Validation , and Write phase.***

i.. Read phase:

- A transaction can *read values of committed data items*.
- However, updates are applied *only to local copies* (versions) of the data items (in database cache).

Concurrency Control Techniques

ii. Validation phase:.

- If the transaction T_i decides that it wants to commit, the DBMS checks whether the transaction could possibly have conflicted with any other concurrently executing transaction.
- While one transaction T_i is being validated, no other transaction can be allowed to commit.
- This phase for T_i checks that, for each transaction T_j that is either committed or is in its validation phase, one of the following conditions holds:

Concurrency Control Techniques

- T_j completes its write phase before T_i starts its read phase.
 - T_i starts its write phase after T_j completes its write phase and the read set of T_i has no item in common with the write set of T_j
 - Both the read_set and write_set of T_i have no items in common with the write_set of T_j .
- When validating T_i , the first condition is checked first for each transaction T_j , since (1) is the simplest condition to check. If (1) is false then (2) is checked and if (2) is false then (3) is checked.
- If none of these conditions holds, the validation fails and T_i is aborted.

iii. Write phase:

- On a successful validation, transactions' updates are applied to the database; otherwise, transactions are restarted.

Concurrency Control Techniques

6. Multiple Granularity Locking

- A *lockable unit* of data defines its *granularity*.
- Granularity can be *coarse* (*entire database*) or it can be *fine* (*an attribute of a relation*).
- Example of *data item granularity*:
 - A field of a database record
 - A database record
 - A disk block/ page
 - An entire file
 - The entire database
- Data item granularity significantly affects *concurrency control performance*.
- Thus, the degree of concurrency is *low for coarse granularity* and *high for fine granularity*.

Concurrency Control Techniques

- Example:
 - A transaction that expects to access most of the pages in a file should probably set a lock on the entire file, rather than locking individual pages or records.
 - If a transaction that requires to access relatively few pages of the file, it is better to lock those pages.
 - Similarly, if a transaction access several records on a page, it should lock the entire page and if it access just a few records, it should lock some those records.
 - This example will hold true, if a lock on the node locks that node and implicitly all its descendants.

Concurrency Control Techniques

- The following diagram illustrates a hierarchy of granularity from *coarse* (*database*) to *fine* (*record*).




Fig 1: Granularity hierarchy

Concurrency Control Techniques

- To manage such hierarchy, in addition to read(S) and write(X), three additional locking modes, called *intention lock modes* are defined:
 - **Intention-shared (IS):** indicates that a shared lock(s) will be requested on some descendent node(s).
 - **Intention-exclusive (IX):** indicates that an exclusive lock(s) will be requested on some descendent node(s).
 - **Shared-intention-exclusive (SIX):** indicates that the current node is locked in shared mode but an exclusive lock(s) will be requested on some descendent nodes(s).
 - If a transaction needs to read an entire file and modify a few of the records in it. i.e it needs S lock on a file & IX lock on some of the contained object


Concurrency Control Techniques

- The set of rules which must be followed for producing serializable schedule are:
 1. The lock compatibility must adhered to.
 2. The root of the tree must be locked first, in any mode.
 3. A node N can be locked by a transaction T in S or IX mode only if the parent node is already locked by T in either IS or IX mode.
 4. A node N can be locked by T in X, IX, or SIX mode only if the parent of N is already locked by T in either IX or SIX mode.
 5. T can lock a node only if it has not unlocked any node (to enforce 2PL policy).
 6. T can unlock a node, N, only if none of the children of N are currently locked by T.

Concurrency Control Techniques

Summery :

Parent locked in	Child can be locked in
IS	IS, S
IX	IS, S, IX, X, SIX
SIX	X, IX
X	None



Concurrency Control Techniques

- To lock a node in *S* mode, a transaction must first lock all its ancestors in *IS* mode. Thus, if a transaction locks a node in *S* mode, no other transaction can have locked any ancestor in *X* mode;
- Similarly, if a transaction locks a node in *X* mode, no other transaction can have locked any ancestor in *S* or *X* mode.
- These two cases ensures that no other transaction holds a lock on an ancestor that conflicts with the requested *S* or *X* lock on the node.

Concurrency Control Techniques

- These locks are applied using the following compatibility matrix:

		Requestor : T1				
		IS	IX	S	SIX	X
Holder: T2	IS	yes	yes	yes	yes	no
	IX	yes	yes	no	no	no
	S	yes	no	yes	no	no
	SIX	yes	no	no	no	no
	X	no	no	no	no	no

- Consider the following three example based Fig 1 diagram given in slide 40
 - T1 wants to update record r111 and r211
 - T2 wants to update all records on page p12
 - T3 wants to read records r11j and the entire f2 file

Concurrency Control Techniques

Multiple Granularity Locking: Lock operation to show a serializable execution

T1	T2	T3
IX(db)		
IX(f1)		
	IX(db)	
		IS(db)
		IS(f1)
		IS(p11)
IX(p11)		
X(r111)		
	IX(f1)	
	X(p12)	
		S(r11j)
IX(f2)		
IX(p21)		
X(r211)		
Unlock (r211)		
Unlock (p21)		
Unlock (f2)		
unlock(db)		

Multiple Granularity Locking: Lock operation to show a serializable execution(continued)

T1	T2	T3
		S(f2)
	unlock(p12)	
	unlock(f1)	
	unlock(db)	
unlock(r111)		
unlock(p11)		
unlock(f1)		
unlock(db)		
		unlock (r111j)
		unlock (p11)
		unlock (f1)
		unlock(f2)

Question & Answer



Thanks !!!